

Aion-BibleQA: Evaluating Retrieval and Citation Faithfulness in Verse-Grounded Bible RAG Systems

Anonymous ACL submission

Abstract

We introduce **Aion-BibleQA**, a 40-question pilot benchmark for evaluating citation faithfulness and false-premise robustness in verse-grounded Bible retrieval-augmented generation (RAG). Existing RAG evaluations measure whether the right document was retrieved, but not whether the system uses retrieved content as citation support in its answer (a gap that matters most when user trust depends on accurate scripture attribution). We build a layered retrieval pipeline that combines exact verse lookup, chapter-aware database queries, numeric keyword suppression, and hybrid semantic ranking, and evaluate it with two metrics: Recall@5 for retrieval coverage and citation_support (an LLM-as-judge score, 0–1) for answer faithfulness. On gold_40 v0.3, our v3 system achieves $R@5 = 0.941$, mean citation_support = 0.978, zero unsupported citations, and a false-premise refusal rate of 1.000 (6/6). To guard against same-family judge bias, a cross-family judge (OpenAI gpt-4.1) re-scores all 40 rows under the same rubric: the two judges agree within one rubric level on every row and report identical unsupported-claim (0.000) and refusal (6/6) rates, so the headline findings are not an artifact of single-judge scoring. Failure analysis shows that the remaining errors trace to retrieval scope (semantic drift in thematic queries and within-chapter verse selection), not to citation misuse, suggesting that retrieval and faithfulness are separable failure modes in Bible RAG systems.

1 Introduction

Bible study applications face a sharper faithfulness challenge than most RAG domains. A user who asks “What does John 3:16 say about love?” expects the system to cite John 3:16, not a topically adjacent verse. If the system retrieves the right chapter but cites the wrong verse within it, the answer sounds plausible while misattributing scripture, a failure invisible to standard retrieval metrics.

Retrieval-augmented generation (Lewis et al., 2020) is the dominant architecture for knowledge-intensive QA. Evaluation typically measures whether retrieval returned relevant documents (Recall@k, MRR) and whether the final answer matches a reference (BLEU, BERTScore, exact match). Neither metric captures whether the system used retrieved content correctly. A system can score $R@5 = 1$ while citing a retrieved passage for a claim it does not support.

We call this gap the **retrieval-faithfulness split**: retrieval correctness and citation faithfulness are distinct failure modes that require separate measurement.

This paper makes four contributions:

- **Aion-BibleQA**, a 40-question benchmark with gold verse annotations across six categories (direct, interpretive, thematic, multi-hop, false-premise, adversarial) designed to stress-test both retrieval and faithfulness.
- A **layered retrieval architecture** (exact verse lookup → chapter-aware DB query → per-chapter coverage guarantee → hybrid semantic re-ranking) achieving $R@5 = 0.941$ on the benchmark.
- **Empirical results** on gold_40 v0.3 using a Gemini LLM-as-judge citation protocol (Table 3): $R@5 = 0.941$, mean citation_support = 0.978, unsupported claim rate = 0.000, false-premise refusal rate = 1.000 (6/6), with a cross-family (GPT) judge corroborating every score within one rubric level (Section 5.4).
- A **failure taxonomy** identifying retrieval scope as the remaining bottleneck: semantic drift and within-chapter verse selection, not citation misuse.

Once retrieval scope is correct, the system cites retrieved verses without misattribution. The split clarifies the engineering agenda: improve what gets retrieved, not how it is cited.

2 Related Work

2.1 Retrieval-Augmented Generation

Lewis et al. (2020) introduced RAG as a parametric–nonparametric hybrid: a dense retrieval component fetches supporting documents, and a seq2seq generator conditions on both the query and retrieved content. Subsequent work improved dense retrieval representations (Karpukhin et al., 2020) and showed that generative models benefit from conditioning on multiple retrieved passages simultaneously (Izacard and Grave, 2021). RAG has been extended to multi-hop settings where answers require aggregating evidence across documents (Tang and Yang, 2024). Retrieval quality caps answer quality: systems cannot generate faithful answers from content they did not retrieve.

In Bible RAG, retrieval quality and answer faithfulness are separable: a system can retrieve the right chapter while citing the wrong verse within it.

2.2 Faithfulness and Grounding Evaluation

Faithfulness (whether an answer is entailed by its supporting context) is distinct from factual correctness (Maynez et al., 2020). RAGAS (Es et al., 2023) decomposes RAG evaluation into faithfulness, answer relevance, and context precision, using an LLM to judge each dimension without ground-truth annotations. G-Eval (Liu et al., 2023) demonstrates strong correlation between GPT-4 judgments and human ratings for faithfulness criteria across summarisation and dialogue tasks. Our Gemini judge scores citation_support; false_premise_refusal adds an orthogonal robustness dimension.

LLM-as-judge suffers from same-family bias: a judge rates outputs from the same model family higher (Zheng et al., 2023). Our judge (Gemini) evaluates answers generated by a Gemini-family model.

2.3 Bible NLP

Resnik et al. (1999) established the Bible as a parallel corpus for cross-lingual alignment, a line extended by the eBible Corpus (Akerman et al., 2023) covering 833 languages. Lima et al. (2025) survey

AI applications to biblical text analysis and identify semantic search and question answering as emerging tasks with limited benchmark coverage. Citation-faithfulness evaluation has not previously been applied to Bible RAG systems.

2.4 False-Premise Robustness

False-premise robustness has been studied in general QA. Hu et al. (2023) construct a dataset of false-premise questions and show that current systems answer false-premise questions as if the premise were true. Bible QA has a specific form of this problem: users frequently misquote or misattribute verses. We include false-premise and adversarial categories in Aion-BibleQA to measure whether the system refuses fabrication requests and corrects misattributions.

3 System Architecture

3.1 Overview

Aion is a verse-grounded Bible QA system deployed as a Supabase Edge Function. Given a user query, the system parses explicit Bible references, retrieves candidate verses, re-ranks them semantically, and generates a grounded answer citing specific verses. Figure 1 shows the pipeline.

3.2 Reference Parsing

A two-pass parser identifies explicit scripture references in the query.

Pass 1 (verse-level): Regex matches BOOK CHAPTER:VERSE patterns (e.g., “John 3:16”, “Ps. 23:1”). An alias map normalises common abbreviations and alternate spellings to canonical book IDs.

Pass 2 (chapter-only): A separate regex captures patterns like “Psalm 23” or “1 Corinthians 15” where no verse is specified. A backtrack fix (lastIndex = match.index + 1 on alias miss) prevents tokens such as “Does 1” from consuming the leading digit of “1 Corinthians”.

Numeric keyword suppression: Parsed chapter references suppress numeric token extraction in the downstream semantic query. Without this, a query for “Psalm 23” generates keyword “23”, retrieving unrelated numerical content (census records, inventory lists).

3.3 Layered Retrieval (v3 Direct-Chapter)

Retrieval proceeds in four layers, ordered by specificity.




Figure 1: Aion v3 retrieval pipeline. Queries with explicit scripture references fork left through exact lookup (L1), chapter-aware DB query (L2), in-chapter semantic re-ranking (L3), and a per-chapter coverage guarantee. Queries without references fork right through hybrid semantic search (L4). Both paths merge at context assembly, which selects the top-5 verses for the answer model.

Layer 1 — Exact verse lookup. For verse-level references, the system queries the `bible_verses` table directly by `(book_id, chapter, verse)`. These rows are added to the candidate set without a semantic ranking step.

Layer 2 — Chapter-aware DB query (`lookupChapterVerses`). For chapter-only references, all verses in the referenced chapter are fetched directly from `bible_verses` by `(book_id, chapter)`. This avoids the v2 design flaw of running an unconstrained semantic search then filtering, which could return semantically close verses from other chapters.

Layer 3 — Semantic re-ranking within chapter (`selectWithinChapters`). Fetched chapter verses are re-ranked by cosine similarity to the query embedding. An IVFFlat index on the embedding column provides approximate nearest-neighbour search. The top- k results are selected.

Per-chapter coverage guarantee. After semantic re-ranking, the system verifies that at least one verse from each referenced chapter appears in the

System	Dataset	R@5	MRR
v1 hybrid-ref	v0.1	0.676	0.552
v2 chapter-ref	v0.2	0.882	0.700
v3 direct-chapter	v0.2	0.882	0.714
v3 direct-chapter	v0.3	0.941	0.773

Table 1: Coarse retrieval ablation across system versions. Each version swaps in one retrieval strategy, isolating the direction of each change rather than individual-layer contribution. The v0.3 gain (+0.059 R@5) is attributable to annotation expansion for two questions (aion_023, aion_033), not architecture change.

final candidate set. If not, it appends the chapter’s first verse. This ensures multi-hop queries referencing two chapters return at least one verse from each.

Layer 4 — Hybrid semantic fallback. For queries with no parsed references, the system runs an unconstrained semantic search over all verses using the query embedding. Keywords extracted from the query bias results toward verses with matching lexical content.

Context assembly. The top-5 verses across all layers are assembled into a context block formatted as `BOOK CHAPTER:VERSE - {text}` on separate lines.

3.4 Answer Generation

We pass the assembled verse context and original query to a Gemini model with a system prompt that instructs it to: answer from the provided verses only; cite specific verses inline; refuse or correct false premises; and decline to fabricate verses or attributions.

3.5 Coarse Retrieval Ablation

The v1→v3 progression doubles as a coarse ablation: each version swaps in one retrieval strategy, so the version-over-version deltas indicate the direction of each design change rather than the isolated contribution of a single layer. v3 matches v2 on the same dataset (R@5 = 0.882; Table 1) despite replacing the retrieval architecture. It fixes aion_035 (Psalm 23 + John 10 multi-hop) but correctly re-breaks aion_036, an accidental success in v2 from an unrestricted semantic fallback.

4 Experiments

4.1 Benchmark: Aion-BibleQA gold_40

Dataset: aion_bibleqa_gold_40_v0.3.jsonl, 40 questions across 6 categories (Table 2).

Category	n	Description
direct	10	Names a specific verse (e.g., “What does John 3:16 say?”)
interpretive	7	Asks for meaning of a named passage
thematic	12	Theme query without naming a verse
multi_hop	5	Requires verses from two different chapters
false_premise	4	Contains a factual error about scripture
adversarial	2	Asks the system to fabricate content

Table 2: Aion-BibleQA category distribution.

Each question has a `gold_verse` (primary expected verse coordinate, e.g., JHN.3.16) and an `acceptable_verse_clusters` list. $R@5=1$ if any retrieved verse matches the gold or any cluster member. We verified all 40 questions against the live `bible_verses` table (BSB translation).

Annotation note: Aion-BibleQA is a pilot benchmark created for this study. Questions cover six failure modes identified during v1 system development, not a random sample of user queries. Results should be interpreted in this context.

4.2 Retrieval Metrics

Recall@5 (R@5): Binary. Score 1 if any of the top-5 retrieved verse coordinates matches `gold_verse` or any cluster member; else 0.

MRR (Mean Reciprocal Rank): $1/r$ where r is the rank of the first matching verse in the ranked list; 0 if no match in top-5.

We compute both metrics per row and report averages per category and overall.

4.3 Citation-Faithfulness Evaluation

Judge model: gemini-3.1-flash-lite.

For each row, the judge receives: (1) the question, (2) the category, (3) expected behaviour from the dataset annotation, (4) the retrieved verse block with full verse texts, and (5) the system answer.

The judge produces a JSON response with two fields:

- `citation_support` (0.0–1.0) for direct/interpretive/thematic/multi_hop rows
- `false_premise_refusal` (0 or 1) for false_premise/adversarial rows

Implementation: We judged all 40 rows in a single pass with three-attempt retry per row and exponential backoff; 0 judge call failures.

Score	Meaning
1.00	Every claim directly grounded in the cited verse text
0.75	Mostly grounded; minor over-reach on one citation
0.50	Some claims go beyond what the verses say
0.25	Most claims loosely connected to the text
0.00	Verses decorative; claims unsupported or fabricated

Table 3: citation_support scoring rubric used by the LLM judge.

4.4 Aggregate Metrics

- **mean_citation_support:** Mean cs across non-refusal rows ($n=34$)
- **unsupported_claim_rate:** Fraction of rows with $cs < 0.5$
- **decorative_citation_rate:** Fraction of rows with $cs \leq 0.25$
- **fp_refusal_rate:** Fraction of false_premise/adversarial rows with $false_premise_refusal = 1$

4.5 Experimental Setup

Bible translation: BSB (Berean Standard Bible) stored in a Supabase `bible_verses` table. Embedding model: text-embedding-3-small (OpenAI; 1,536 dimensions), stored as `halfvec(1536)` in `pgvector` with an `IVFFlat` index (`halfvec_cosine_ops`). Answer model: Gemini (same model family as the primary judge; see the Limitations section), deployed as a Supabase Edge Function on Deno Deploy infrastructure. Judges: gemini-3.1-flash-lite (primary) and OpenAI gpt-4.1 (secondary, cross-family robustness check; Section 5.4). Both receive an identical rubric and prompt at temperature 0.1. Benchmark runner: calls the live Edge Function over HTTPS with SSE streaming from a consumer MacBook. All benchmark runs are frozen as JSONL artefacts; results are reproducible from those frozen files.

5 Results

5.1 Main Results

v3 on gold_40 v0.3 achieves $R@5=0.941$, `citation_support`=0.978, and `fp_refusal`=1.000 (Table 4), suggesting that once retrieval scope is correct, the system cites retrieved verses faithfully.

Metric	Score
R@5	0.941
MRR	0.773
Mean citation_support	0.978
Unsupported claim rate ($cs < 0.5$)	0.000
Decorative citation rate ($cs \leq 0.25$)	0.000
False-premise / adversarial refusal	1.000 (6/6)

Table 4: v3 system on gold_40 v0.3: overall metrics.

Category	n	R@5	MRR	cs	fp
direct	10	1.000	1.000	1.000	—
interpretive	7	1.000	1.000	1.000	—
thematic	12	0.917	0.667	0.979	—
multi_hop	5	0.800	0.600	0.900	—
false_premise	4	—	—	—	1.000
adversarial	2	—	—	—	1.000

Table 5: Per-category results. cs=citation_support; fp=fp_refusal. 95% Wilson CIs for R@5: overall [0.80, 0.97], multi_hop [0.38, 0.96].

5.2 Per-Category Breakdown

Direct and interpretive categories. R@5 = 1.00 and citation_support = 1.00 for all 17 questions (Table 5). The exact verse lookup layer handles all named verse queries; semantic re-ranking does not introduce wrong verses when the reference is explicit.

Thematic and multi-hop contain the remaining retrieval failures. Two questions fail R@5: aion_027 (thematic, grace semantic drift) and aion_036 (multi-hop, IVFFlat boundary). Citation faithfulness remains high for both categories ($cs = 0.979$ and 0.900), showing that the system cites retrieved content correctly even when retrieval returned suboptimal verses.

Refusal categories. All six refusal-category questions scored 1 (4 false-premise, 2 adversarial). No fabricated content appeared in any answer. The adversarial sample ($n = 2$) is too small for confident generalization; these results are directionally encouraging.

5.3 Sub-1.0 Rows

Two rows scored below citation_support = 1.0.

aion_021 (thematic, $cs = 0.75$): EPH.6.19 cited for a gratitude context. The verse records Paul requesting prayer for his ministry, not a general statement about thankfulness. A minor over-reach.

aion_035 (multi-hop, $cs = 0.50$): JHN.10.1 cited instead of JHN.10.11. The per-chapter coverage guarantee appends the first verse of JHN.10 (the parable of the gate) rather than the theologically

Metric	Gemini	GPT
Mean citation_support ($n = 34$)	0.978	0.941
Unsupported rate ($cs < 0.5$)	0.000	0.000
Decorative rate ($cs \leq 0.25$)	0.000	0.000
False-premise refusal	1.000	1.000

Table 6: Two-judge panel on gold_40 v0.3. The cross-family judge reproduces the unsupported-claim and refusal rates exactly; mean citation_support differs by 0.037.

central verse (“I am the good shepherd”). Notably, R@5 = 1 for this question (PSA.23.1 retrieved correctly); the failure is within-chapter verse selection, not retrieval coverage. aion_035 achieves R@5 = 1 yet $cs = 0.50$, a direct illustration of the retrieval-faithfulness split.

5.4 Multi-Judge Robustness

The primary judge (Gemini) shares a model family with the answer model, which risks same-family leniency (Zheng et al., 2023). To test whether the headline scores survive a change of judge, we re-scored all 40 rows with a cross-family judge (OpenAI gpt-4.1) under the identical rubric and prompt (Table 6).

The two judges agree exactly on 31/40 rows and *within one rubric level on all 40* (every disagreement is a single 0.25 step). Critically, the conclusions that do not depend on a fractional score are judge-invariant: both judges place zero rows below the unsupported threshold and confirm all 6 refusals. The nine disagreements are minor over-reach calls (e.g. gpt-4.1 rates aion_021 at 1.00 where Gemini gives 0.75, and the reverse on aion_027); none crosses a decision boundary. We report this as an automated multi-model robustness check, not as expert human adjudication.

6 Discussion

6.1 The Retrieval-Faithfulness Split

Retrieval correctness and citation faithfulness are separable. The two metrics diverge in multi_hop: R@5 = 0.800 while citation_support = 0.900. A system can retrieve the wrong verse (R@5 = 0) yet cite whatever it retrieved correctly ($cs = 1.00$): aion_027 illustrates this. Conversely, a system can achieve R@5 = 1 while citing the wrong verse within a chapter ($cs = 0.50$): aion_035 illustrates this.

Both metrics are necessary: high R@5 does not prevent citation misuse, and high citation_support

Class	Example	Fix
Semantic drift	aion_027: “grace” → salutation formulas	Query expansion salutation penalty (v3.1)
IVFFlat boundary	aion_036: PHP.4.6 in-consistently absent	Per-chapter vector RPC (v4)
Per-chapter wrong verse	aion_035: JHN.10.1 vs JHN.10.11	Per-chapter vector RPC (v4)

Table 7: Remaining failure taxonomy. None is a citation misuse failure.

does not confirm that retrieved content was relevant.

6.2 Failure Taxonomy

All remaining failures fall into one of three classes (Table 7).

In every case where retrieval returned relevant content, the answer model cited it correctly. Fixing retrieval scope addresses both remaining failure classes.

6.3 The aion_027 Grace Drift

“Grace” in Pauline letters has two semantic neighbourhoods: salvific grace (EPH.2.8–9: “For by grace you have been saved through faith”) and epistolary grace (greeting formulas of the form “Grace and peace to you. . .”). The epistolary cluster dominates the current semantic search surface: these short, formulaic verses have embeddings closer to the unqualified query “grace” than the longer salvific passage.

Two candidate fixes: (1) query expansion, prepending “saved by” or “salvation” to thematic grace queries at retrieval time; (2) salutation suppression, a post-retrieval filter penalising verses matching the pattern “Grace and peace to you from X”. The second is more targeted; the first generalises more broadly to other semantic drift cases.

6.4 v4 Architecture Roadmap

aion_035 and aion_036 share a root cause: the per-chapter guarantee and IVFFlat search operate at chapter granularity, not verse-within-chapter granularity.

The v4 fix is a search_verses_in_chapter Supabase RPC running constrained vector search within a single chapter, parameterised by query embedding, book ID, chapter number, and return count k . For JHN.10, a query embedding for “God as shepherd” would rank JHN.10.11 above

JHN.10.1. The per-chapter guarantee becomes the top result of this per-chapter search rather than the chapter’s first verse. This is a well-defined engineering task.

7 Conclusion

We introduced Aion-BibleQA, a 40-question benchmark for citation-faithfulness and false-premise robustness in verse-grounded Bible RAG. Our v3 retrieval system, combining exact verse lookup, direct chapter-aware database queries, and hybrid semantic re-ranking, achieves $R@5 = 0.941$ on gold_40 v0.3. A Gemini LLM-as-judge evaluation shows mean citation_support = 0.978, zero unsupported citations, and a 1.000 false-premise refusal rate.

Retrieval and citation faithfulness are distinct failure modes. Once the right verses reach the context window, the system cites them correctly. The remaining errors (semantic drift in thematic queries, IVFFlat boundary effects in multi-hop retrieval, and within-chapter verse selection) are retrieval scope problems.

Future priorities:

- Expert (theological) human annotation, beyond the automated cross-family judge agreement already reported.
- Expansion to 200+ questions sampled from real users.
- Thematic query expansion for grace-like semantic drift (v3.1).
- Per-chapter vector search RPC for within-chapter verse selection (v4).

We release Aion-BibleQA, the retrieval system, the judge harness, and all benchmark artefacts alongside this paper.

Limitations

Automated judging, no expert human validation. Citation-faithfulness scores come from LLM judges, and the primary judge (Gemini) shares a model family with the answer model, which risks same-family leniency (Zheng et al., 2023). We mitigate this with a cross-family judge (gpt-4.1) on all 40 rows: the two judges agree within one rubric level on every row and report identical unsupported-claim and refusal rates (Section 5.4), so the headline conclusions are robust to the choice of judge.

We do not, however, claim expert theological validation: we evaluate only whether generated claims are supported by the cited verse text, under a reproducible multi-model protocol. Expert human annotation, and a larger automated panel (a third judge harness is included but unrun for lack of an API key), remain future work.

Pilot benchmark size. gold_40 v0.3 contains 40 questions. Statistical uncertainty is high: R@5 is averaged over the 34 non-refusal rows, so one question shifts it by roughly 3% (1 of 34). Results in the multi_hop category (n=5) and adversarial category (n=2) are especially noisy. A 95% CI on multi_hop R@5=0.800 with n=5 spans roughly ± 0.30 .

Constructed benchmark, not user-sampled. Questions target known failure modes from v1 development, not the distribution of queries real users ask. The benchmark was created by the same team that built the system under evaluation; scores may overstate performance relative to an independently assembled evaluation set. Real-world R@5 may differ.

Single translation (BSB). All verses come from the Berean Standard Bible. Results may not generalise to systems using KJV, NIV, ESV, or other translations, where verse wording differs and embeddings shift.

IVFFlat non-determinism. Across three repeated runs on gold_40 v0.2, three questions showed R@5 variance (aion_006, aion_007, aion_036). Run-to-run R@5 ranged from 0.853 to 0.882 on the same question set. Benchmark runs in this paper were each performed once and frozen as canonical JSONL artefacts; results are reproducible from those frozen files.

Coarse, not controlled, ablation. The v1→v2→v3 progression (Table 1) acts as a coarse ablation showing the direction of each design change, but each version alters more than one component at once, so the isolated contribution of an individual retrieval layer is not quantified.

Acknowledgments

We thank the Berean Bible translators for making the BSB publicly available under a generous licence, which made this corpus-based study possible.

Ethics Statement

This study uses the Berean Standard Bible (BSB), a publicly available translation released under a licence permitting research use. No personal data were collected; Aion-BibleQA contains no human subject information. The judge and answer models are commercial APIs; no model weights were modified or redistributed. The benchmark targets a low-risk religious reference task with no plausible pathway to harm from retrieval failures.

References

- Vesa Akerman, David Baines, Damian Daspit, and Ulf Hermjakob. 2023. The eBible corpus: Data and model benchmarks for Bible translation for low-resource languages. ArXiv:2304.09919.
- Shahul Es, Jithin James, Luis Espinosa-Anke, and Steven Schockaert. 2023. RAGAS: Automated evaluation of retrieval augmented generation. ArXiv:2309.15217.
- Shengding Hu, Yulong Luo, Huadong Wang, Xingyi Cheng, and Zhiyuan Liu. 2023. Won’t get fooled again: Answering questions with false premises. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics*. ACL 2023.
- Gautier Izacard and Edouard Grave. 2021. Leveraging passage retrieval with generative models for open domain question answering. In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics*. EACL 2021. arXiv:2007.01282.
- Vladimir Karpukhin, Barlas Oğuz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. 2020. Dense passage retrieval for open-domain question answering. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*. EMNLP 2020. arXiv:2004.04906.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems*. ArXiv:2005.11401.
- B. C. Lima, N. Omar, I. Avansi, and L. N. de Castro. 2025. Artificial intelligence applied to the analysis of biblical scriptures: A systematic review. *Analytics*, 4(2):13.
- Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. 2023. G-Eval: NLG evaluation using GPT-4 with better human alignment. In *Proceedings of the 2023 Conference on*

Empirical Methods in Natural Language Processing.
EMNLP 2023. arXiv:2303.16634.

Joshua Maynez, Shashi Narayan, Bernd Bohnet, and
Ryan McDonald. 2020. On faithfulness and fac-
tuality in abstractive summarization. In *Proceed-
ings of the 58th Annual Meeting of the Associa-
tion for Computational Linguistics.* ACL 2020.
arXiv:2005.00661.

Philip Resnik, Mari Broman Olsen, and Mona Diab.
1999. The Bible as a parallel corpus: Annotating
the ‘Book of 2000 tongues’. *Computers and the
Humanities*, 33(1–2).

Yixuan Tang and Yi Yang. 2024. MultiHop-RAG:
Benchmarking retrieval-augmented generation for
multi-hop queries. ArXiv:2401.15391.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan
Zhuang, Zhonghao Wu, Yonghao Zhuang, Zi Lin,
Zhuohan Li, Dacheng Li, Eric Xing, Hao Zhang,
Joseph E. Gonzalez, and Ion Stoica. 2023. Judging
LLM-as-a-judge with MT-Bench and Chatbot Arena.
In *Advances in Neural Information Processing Sys-
tems.* NeurIPS 2023 Datasets and Benchmarks Track.
arXiv:2306.05685.